Poseidon

The Trusted AI-Native Money Platform

Group 7 | SF Shinji Fujiwara | SB Sean Beecroft | MH Michael Hinckley | AK Arun Kumar

We use AI tools to assist development. We review, test, and take responsibility for the final output.

STRUCTURAL ISSUE

The Coordination Gap

Your financial data is fragmented. You are responsible for coordination.

Banking
MANUAL
Cross-institution fund move

Credit
MANUAL
Static auto-pay

Investment
MANUAL
Isolated from cash flow context

Budget
MANUAL
Tracks but can't act on funds

You are the integration layer.

$133/mo
Subscription waste
Per active user per month (C+R 2024).

$12.5B/yr
Fraud losses
Annual fraud and theft across US (FTC, 2024)

$5B/yr
Overdraft fees
Annual overdraft & Non Sufficient Fund fee charged. (CFPB, 2023)

Fintech(intuit mint , etc) solved visibility. What can we solve next?

WHY NOW?

Environmental Shift

Three forces converging to make trustworthy personal finance AI.

Open Banking

- APIs are standardizing
- Data processing rules clarified(GDPR, etc)

API Regulation Data

AI-Native Expectation

Users now expect proactive, personalized financial guidance

Behavioral Shift AI Native Expectation

2021 2023 2025

AI Economics

Inference cost dropping 10x annually (Epoch AI, 2025)

Infrastructure + economics + user demand converged in 2025. Ready for scale.

GOVERNANCE BY DESIGN ARCHITECTURE

Poseidon: 4 Engines

Protect, Grow, Execute, and Govern as one auditable system.

Govern

Ensures auditability of all engines

- Compliance
- AI Governance
- Full Auditability
- Transparent UI

Protect

Personalized ML models across your accounts

- Detection
- Fraud
- Subscription
- Anomaly

Grow

Short & long term recommendation

- Cash Flow Forecast
- Portfolio Analysis
- Optimization
- Actionable insights

Execute

Human approval based automated execution

- Cross-Engine Link
- Human Approval
- Reversible Actions
- Centralized UI

ARCHITECTURE PRINCIPLE

Deterministic models compute | GenAI explains | AI Agents execute | Humans confidently approve

DIFFERENTIATOR

Beyond Aggregation

From dashboards to prediction, and approval-first execution.

COMPETITIVE MOAT DEEPENS →

COMMODITIZED
Everyone does this

✓ **Aggregation**
Multi-bank account linking

✓ **Budgeting**
Rule-based spend tracking

EMERGING
Some are trying

△ **Dashboard**
Visibility without action

△ **AI-Powered Insights**
Limited

ONLY POSEIDON
No one else

★ **Predictive Intelligence**
Personalized ML models

★ **Explainable AI**
Plain English explanation with low temperature + contributing factors

★ **Approval-first Execution**
Human-in-the-loop automation

ROADMAP

Compliance-first roadmap

Phased execution plan with measurable progress.

Compliance

Phase 1 0-3 Months

Foundation Establish compliant foundation.

- **Compliance**
  Bank-grade protocols, SOC2, privacy-by-design
- **Governance**
  AI ethics board, risk assessment framework
- **LLMOps / MLOps**
  Model lifecycle management, monitoring
- **AI explainability**
  Transparency-by-design
- **Detection and Grow pilots**
  Customer data integration, controlled testing

Phase 2 3-12 Months

Automation

Frontier POC execution engine

Precision ⩾ 70% | Reverse Option | Workflow & dashboard

Phase 3 12-15 Months

Stability

Break-even Prove reliability and reach break-even economics.

Users ~180K | Precision ⩾ 80% | Availability ⩾ 99.9%

Phase 4 15+ Months

More Data Source

Scale Increase user base, data coverage, and ML model scope.

Users ~251K | Precision ⩾ 90% | FP ⩽ 5%

Phase 5 | B2B, White Label

Distribution

DEMO

Product Walkthrough

See Poseidon in action — from alert to execution

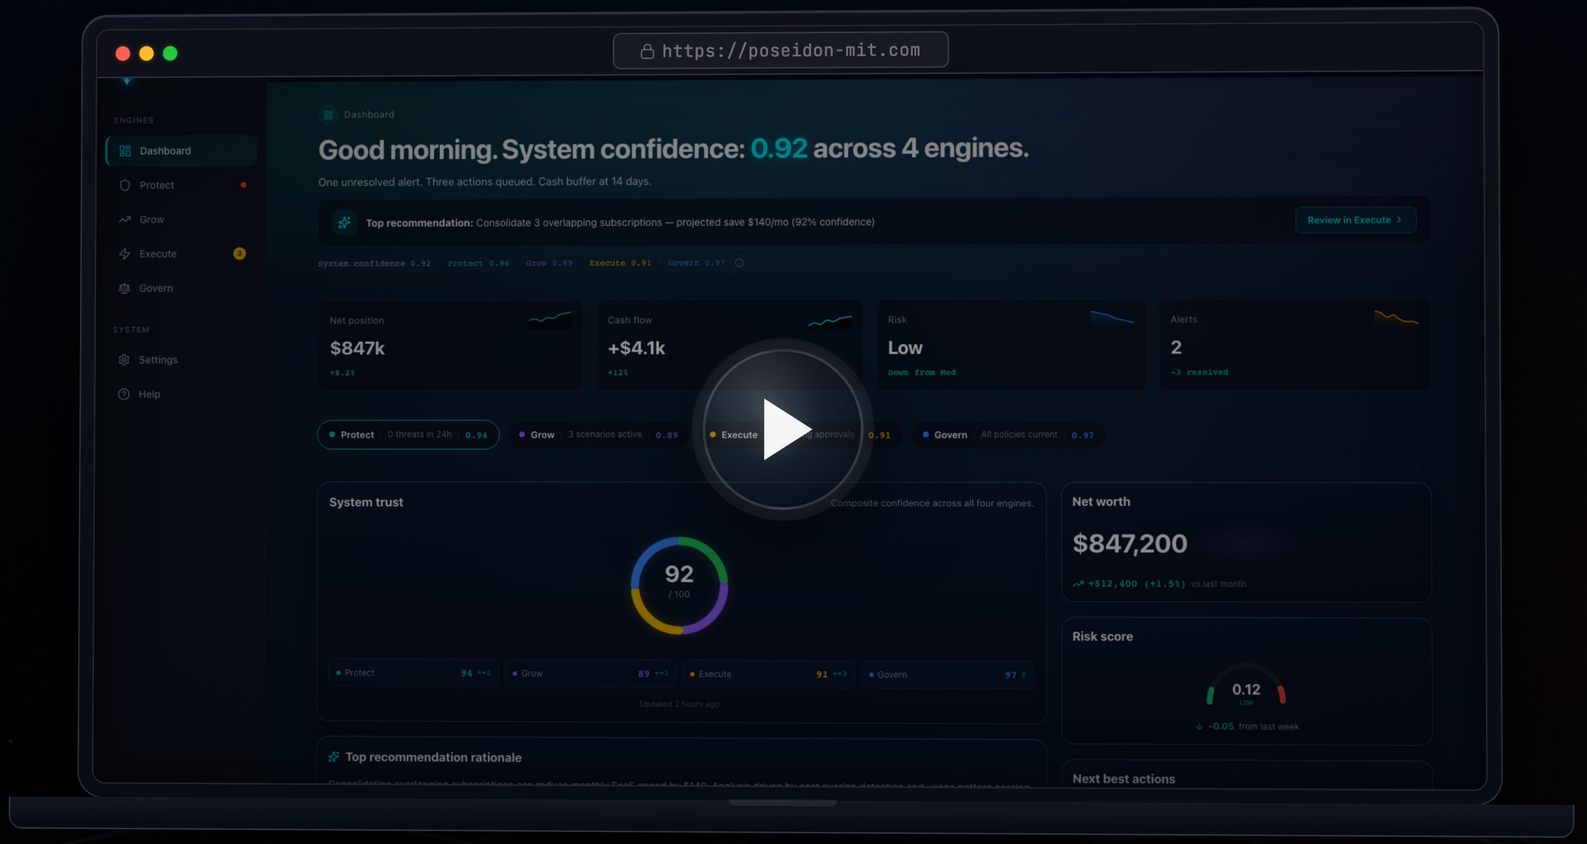

BUSINESS MODEL

AI-powered profitable growth

Efficient operating model powered by AI

UNIT ECONOMICS (AT SCALE)

4.3X	77%	Month 12
Value / Cost for customer	Profit Margin	Operating Breakeven

MARKET OPPORTUNITY

$3.5B	$1.5B	$46M
Total Market	Target Segment	3-Year Capture

Conservatively calculated from Federal Reserve (2024), Fortune Business Insights (2025), Allied Market Research (2025)

PRICING PLANS

FREE	Dashboard + Govern + Basic Protect	$0 Starting point
PLUS	Full Protect + Grow + Execute (limited)	$7.99 per month 3.5% of free users upgrade
PRO	Full capability	$19.99 per month 0.8% choose Pro + upgrades

CUMULATIVE PROFIT ($MILLION)

STRESS SCENARIO

SCENARIO	PAID CONV.	CHURN	BREAKEVEN	PAYBACK
Base	3.5%	5%	M12	M19
Conservative	2.0%	7%	M16	M26
Pessimistic	1.5%	10%	M18	M31

SUMMARY

Poseidon Strategy Summary

VISION

Establish the trusted financial platform where AI coordination serves human financial wellbeing

Governance first

Meet regulatory expectation, provide further value

- Regulatory Compliance
- ML/LLMOps
- Explainable AI
- Reversible AI

Architecture

Unified AI Architecture

Deterministic models compute	GenAI explains	AI Agents execute	Humans confidently approve
ML models calculate with precision	Plain English explanation	Workflow orchestration	Human-in-the-loop with centralized UI

Business Model

Sustainable business, measurable progress Appendix ›

77%	Month 12	4.3X
Gross Margin	Op. Break-even	Value / Cost for customer

ONE YEAR REFLECTION

Tough — and totally worth it.

Try the prototype

https://poseidon-mit.com

SF Shinji Fujiwara

SB Sean Beecroft

MH Michael Hinckley

AK Arun Kumar

Poseidon | MIT Professional Education

Appendix

RISK REGISTER

Inherent Risk and Mitigation

Six core risks, key controls, and residual exposure at a glance.

#1 Impact: High Likelihood: High Residual: Moderate

Model drift from behavior and fraud-pattern shifts.

MITIGATION

- Continuous precision/recall and drift monitoring.
- Scheduled and threshold-triggered retraining.

#2 Impact: High Likelihood: High Residual: Moderate

External, insider, and supply-chain security threats.

MITIGATION

- Defense-in-depth: least privilege, MFA, segmentation, SIEM/SOC.
- Secure SDLC plus tested incident-response playbooks.

#3 Impact: Medium Likelihood: High Residual: Moderate

LLM hallucination or misleading explanations.

MITIGATION

- Deterministic services compute all numeric outputs.
- Grounded responses with post-generation validation and rollback controls.

#4 Impact: High Likelihood: Medium Residual: Moderate

Bias, fairness, and disparate-impact risk.

MITIGATION

- Fairness testing and continuous output monitoring.
- Pre-deployment harm evaluation and feedback remediation loop.

#5 Impact: High Likelihood: Medium Residual: Moderate

Regulatory, privacy, and data-protection exposure.

MITIGATION

- Data residency, minimization, and PII sanitization before LLM use.
- Encryption, key rotation, IAM, and periodic control audits.

#6 Impact: High Likelihood: Medium Residual: Low

System resiliency and availability degradation.

MITIGATION

- Auto-healing and automatic failover.
- DR with RTO/RPO tests and BCP non-LLM fallback.